

INT304 Support Vector Machine (SVM)

The **Support Vector Machine (SVM)** is a powerful supervised learning model that finds the optimal hyperplane separating data from different classes with the **maximum margin**.

SVM works well in high-dimensional spaces and is effective when the number of features exceeds the number of samples.

Objective Function

For linearly separable data, SVM solves:

$$\min_{w,b} \quad \frac{1}{2} \|w\|^2 \quad \text{subject to} \quad y_i(w^T x_i + b) \geq 1$$

This maximizes the margin between classes while ensuring correct classification.

Kernels

To handle non-linearly separable data, SVM uses **kernel functions** to map input data into a higher-dimensional space. Common kernels:

- 'linear' : $K(x_i, x_j) = x_i^T x_j$
- 'poly' : $K(x_i, x_j) = (\gamma x_i^T x_j + r)^d$
- 'rbf' : $K(x_i, x_j) = \exp(-\gamma \|x_i - x_j\|^2)$
- 'sigmoid' : $K(x_i, x_j) = \tanh(\gamma x_i^T x_j + r)$

We will use `sklearn`'s `SVC` (Support Vector Classifier) to implement SVM efficiently.

`sklearn.svm.SVC`

```
In [23]: # import necessary libs
from sklearn.datasets import load_digits
from sklearn import svm, metrics
from sklearn.model_selection import train_test_split
import numpy as np
import matplotlib.pyplot as plt
```


```
In [24]: # Load dataset
digits = load_digits()
X, y = digits.data, digits.target

# Show first few samples
fig, axes = plt.subplots(1, 5, figsize=(10, 2))
for i, ax in enumerate(axes):
    ax.imshow(digits.images[i], cmap='gray')
    ax.set_title(f"Label: {digits.target[i]}")
    ax.axis('off')
plt.tight_layout()
plt.show()

# Split into train/test
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=1/7, random_state=42)

# Train SVM
clf = svm.SVC(kernel='linear') # You may want to tune kernel, C, gamma
clf.fit(X_train, y_train)

# Predict & Evaluate
y_pred = clf.predict(X_test)
print("Accuracy:", metrics.accuracy_score(y_test, y_pred))
```



Accuracy: 0.9766536964980544

Parameters you may tune for training

- **C : float, default = 1.0**

Regularization parameter. The strength of the regularization is inversely proportional to C . Must be strictly positive. The penalty is a squared L_2 penalty.

- **kernel : 'linear' , 'poly' , 'rbf' , 'sigmoid' , 'precomputed' or callable, default = 'rbf'**

Specifies the kernel type to be used in the algorithm.

If none is given, 'rbf' will be used.

If a callable is given, it is used to pre-compute the kernel matrix from data matrices; that matrix should be an array of shape $(n_samples, n_samples)$.

For an intuitive visualization of different kernel types, see *Plot classification boundaries with different SVM Kernels*.

- **degree : int, default = 3**

Degree of the polynomial kernel function ('poly'). Must be non-negative.

Ignored by all other kernels.

- **coef0 : float, default = 0.0**

Independent term in kernel function. It is only significant in 'poly' and 'sigmoid'.

- **tol : float, default = $1e-3$**

Tolerance for stopping criterion.

Observe:

- Does overfitting or underfitting occur with certain parameter choices?

`sklearn.svm.SVR`

Support Vector Regression (SVR) is an extension of Support Vector Machines (SVM) to regression tasks. Unlike linear regression, which minimizes squared errors, SVR aims to fit a function that deviates from actual targets by at most ϵ , while also being as flat as possible.

Optimization Objective

SVR minimizes the following:

$$\min_{w,b,\xi,\xi^*} \frac{1}{2} \|w\|^2 + C \sum_{i=1}^n (\xi_i + \xi_i^*)$$

Subject to:

$$\begin{aligned} y_i - (w^T x_i + b) &\leq \epsilon + \xi_i \\ (w^T x_i + b) - y_i &\leq \epsilon + \xi_i^* \\ \xi_i, \xi_i^* &\geq 0 \end{aligned}$$

Where:

- C is the regularization parameter
- ϵ defines a margin of tolerance (called the ϵ -tube)
- ξ_i, ξ_i^* are slack variables for errors outside the ϵ -tube

SVR works well when the data is noisy and you want to balance between **model complexity** and **precision tolerance**.

```
In [25]: # import necessary libs
import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.svm import SVR
from sklearn.metrics import mean_squared_error, mean_absolute_error
from sklearn.preprocessing import MinMaxScaler, StandardScaler
import matplotlib.pyplot as plt
```

```
In [26]: data_url = "http://lib.stat.cmu.edu/datasets/boston"
raw_df = pd.read_csv(data_url, sep="\s+", skiprows=22, header=None)
```

```

X = np.hstack([raw_df.values[::2, :], raw_df.values[1::2, :2]])
Y = raw_df.values[1::2, 2]
x_train, x_test, y_train, y_test = train_test_split(X, Y, test_size=0.2, random_state=42)
scaler = MinMaxScaler()
norm_x_train = scaler.fit_transform(x_train)
norm_x_test = scaler.transform(x_test)

model = SVR() # you need tuning here
model.fit(norm_x_train, y_train)
y_pred = model.predict(norm_x_test)
test_mse = mean_squared_error(y_test, y_pred)
test_mae = mean_absolute_error(y_test, y_pred)
print(f'Test MSE: {test_mse}, Test MAE: {test_mae}')

```

Test MSE: 27.907001783282002, Test MAE: 2.94557507610235


Plot actual vs predicted values. Are there consistent under- or over-estimations?

```

In [27]: # plot actual vs predicted values
import matplotlib.pyplot as plt

plt.scatter(y_test, y_pred, alpha=0.7)
plt.plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()], 'r--')
plt.xlabel("Actual")
plt.ylabel("Predicted")
plt.title("SVR Prediction vs Actual")
plt.grid(True)
plt.show()

```



🔧 Parameters you may tune for training

- **kernel** : 'linear' , 'poly' , 'rbf' , 'sigmoid' , 'precomputed' or callable, default = 'rbf'
Specifies the kernel type to be used in the algorithm.
If none is given, 'rbf' will be used.
If a callable is given, it is used to precompute the kernel matrix.
- **degree** : int, default = 3
Degree of the polynomial kernel function ('poly'). Must be non-negative.
Ignored by all other kernels.
- **coef0** : float, default = 0.0
Independent term in kernel function. It is only significant in 'poly' and 'sigmoid' .
- **tol** : float, default = $1e-3$
Tolerance for stopping criterion.
- **C** : float, default = 1.0
Regularization parameter. The strength of the regularization is inversely proportional to C .

Must be strictly positive. The penalty is a squared L_2 penalty.

- **epsilon : float, default = 0.1**

Epsilon in the epsilon-SVR model.

It specifies the epsilon-tube within which no penalty is associated in the training loss function with points predicted within a distance epsilon from the actual value.

Must be non-negative.

In []: